

Q3 Cat and Mouse Game

a) Stationary Distributions

Stationary Distribution of Cat Chain :

Let $\pi_1 = P(X_t = \text{Room 1})$ and $\pi_2 = P(X_t = \text{Room 2})$, then:

$$\pi_1 = P(X_{t-1} = \text{Room 1}) \cdot P(X_t = \text{Room 1} \mid X_{t-1} = \text{Room 1}) + P(X_{t-1} = \text{Room 2}) \cdot P(X_t = \text{Room 1} \mid X_{t-1} = \text{Room 2})$$

$$\pi_2 = P(X_{t-1} = \text{Room 1}) \cdot P(X_t = \text{Room 2} \mid X_{t-1} = \text{Room 1}) + P(X_{t-1} = \text{Room 2}) \cdot P(X_t = \text{Room 2} \mid X_{t-1} = \text{Room 2})$$

Using the transition probabilities:

- $P(X_t = \text{Room 1} \mid X_{t-1} = \text{Room 1}) = 0.2$
- $P(X_t = \text{Room 1} \mid X_{t-1} = \text{Room 2}) = 0.8$
- $P(X_t = \text{Room 2} \mid X_{t-1} = \text{Room 1}) = 0.8$
- $P(X_t = \text{Room 2} \mid X_{t-1} = \text{Room 2}) = 0.2$

With $\pi_1 + \pi_2 = 1$, we get:

$$\pi_1 = \pi_1 \cdot 0.2 + \pi_2 \cdot 0.8$$

$$\pi_2 = \pi_1 \cdot 0.8 + \pi_2 \cdot 0.2$$

Final Stationary Probabilities:

- $\pi_1 = 0.5$
- $\pi_2 = 0.5$

Stationary Distribution of Mouse Chain :

Let $\mu_1 = P(Y_t = \text{Room 1})$ and $\mu_2 = P(Y_t = \text{Room 2})$, then:

$$\mu_1 = P(Y_{t-1} = \text{Room 1}) \cdot P(Y_t = \text{Room 1} \mid Y_{t-1} = \text{Room 1}) + P(Y_{t-1} = \text{Room 2}) \cdot P(Y_t = \text{Room 1} \mid Y_{t-1} = \text{Room 2})$$

$$\mu_2 = P(Y_{t-1} = \text{Room 1}) \cdot P(Y_t = \text{Room 2} \mid Y_{t-1} = \text{Room 1}) + P(Y_{t-1} = \text{Room 2}) \cdot P(Y_t = \text{Room 2} \mid Y_{t-1} = \text{Room 2})$$

Using the transition probabilities:

- $P(Y_t = \text{Room 1} \mid Y_{t-1} = \text{Room 1}) = 0.7$
- $P(Y_t = \text{Room 1} \mid Y_{t-1} = \text{Room 2}) = 0.6$
- $P(Y_t = \text{Room 2} \mid Y_{t-1} = \text{Room 1}) = 0.3$
- $P(Y_t = \text{Room 2} \mid Y_{t-1} = \text{Room 2}) = 0.4$

With $\mu_1 + \mu_2 = 1$, we get:

$$\mu_1 = \mu_1 \cdot 0.7 + \mu_2 \cdot 0.6$$

$$\mu_2 = \mu_1 \cdot 0.3 + \mu_2 \cdot 0.4$$

Final Stationary Probabilities:

- $\mu_1 = \frac{2}{3}$
- $\mu_2 = \frac{1}{3}$

b) Is it a markov chain?

Yes, the sequence Z_0, Z_1, Z_2, \dots **is** a Markov chain.

In this case:

- The cat and mouse move independently.
- Each next position depends only on the **current position**, not on the history.
- Hence, the joint transition probability:

$$P(Z_{n+1} \mid Z_n) = P(\text{Cat}_{n+1} \mid \text{Cat}_n) \cdot P(\text{Mouse}_{n+1} \mid \text{Mouse}_n)$$

depends **only on** Z_n .

Therefore, $\{Z_n\}$ satisfies the Markov property and is a **valid Markov chain** with 4 states.

Q5 Stock Price Model

a) Is the Stock Price Recurrent?

No, it is not recurrent.

- The state space is **infinite** in both directions (prices can go very high in theory)
- Because of the average positive drift, the stock tends to go up over time

b) Does the stationary distribution of the stock price exist?

There is **no** stationary distribution because the stock price has a positive drift and tends to increase indefinitely over time. This means the process does not settle into a long-term equilibrium and does not return to states often enough.

c) Does the stock price reach 130 INR before 1:00 PM?

We will use dynamic programming with mpmath for high precision.

```
In [1]: from mpmath import mp, mpf
```

```
# Set decimal precision  
mp.dps = 10
```

```
tick_size = 0.01  
start_price = 120.00  
target_price = 130.00  
total_steps = 3*60*60//5
```

```
prob_up = mpf('0.1')  
prob_same = mpf('0.85')
```

```

prob_down = mpf('0.05')

# Create DP table p[t][s] where: t is the step number and s is the price index
p = [[mpf('0') for _ in range(2 * total_steps + 1)] for _ in range(total_steps + 1)]
# Initialize starting price index
p[0][total_steps] = mpf('1.0')

# Fill DP table
for step in range(1, total_steps + 1):
    for price in range(2 * total_steps + 1):
        if p[step - 1][price] != mpf('0'):
            # Non zero probability from previous step will propagate
            # to the next step based on the probabilities
            p[step][price - 1] += p[step - 1][price] * prob_down
            p[step][price] += p[step - 1][price] * prob_same
            p[step][price + 1] += p[step - 1][price] * prob_up

# Sum up final probabilities for prices >= target
index_shift = int((target_price - start_price) / tick_size)
final_probs = p[total_steps][total_steps + index_shift:]
prob_hit = sum(final_probs)
print(f"Estimated probability using high-precision DP: {prob_hit}")

```

Estimated probability using high-precision DP: 1.069032254e-432

We can also use log-space probabilities for precision, but that was found to be more computationally expensive than mpmath.